

SILENT VISION

A fixed phrase. A silent clip. An inspectable decision.

Personalized visual command recognition - not open-vocabulary lipreading.


FOUR-PHRASE BILINGUAL CATALOG

你好，请帮我打开灯 | 你吃饭了吗？


Hello, please turn on the light. | Have you eaten?

The user selects the language before recording. Selected-language softmax scores only that language, then exact catalog text or heuristic UNKNOWN is returned.

HOW IT WORKS



WHERE IT FITS



SAFETY RULE

Both calibrated gates must pass. Heuristic UNKNOWN carries no exact phrase text and cannot execute.

AMD RADEON + ROCM

CPU handles video and face preprocessing. ROCm PyTorch runs the learned fixed-phrase model on Radeon cuda:0.

EVIDENCE STATUS

Pending: English recordings, bilingual training, ROCm execution, and final evaluation. No bilingual accuracy claim before the final report. Small-data smoke proves execution only.



SOURCE REPOSITORY