

SILENT VISION

A fixed phrase. A silent clip. An inspectable decision.

Personalized visual command recognition - not open-vocabulary lipreading.

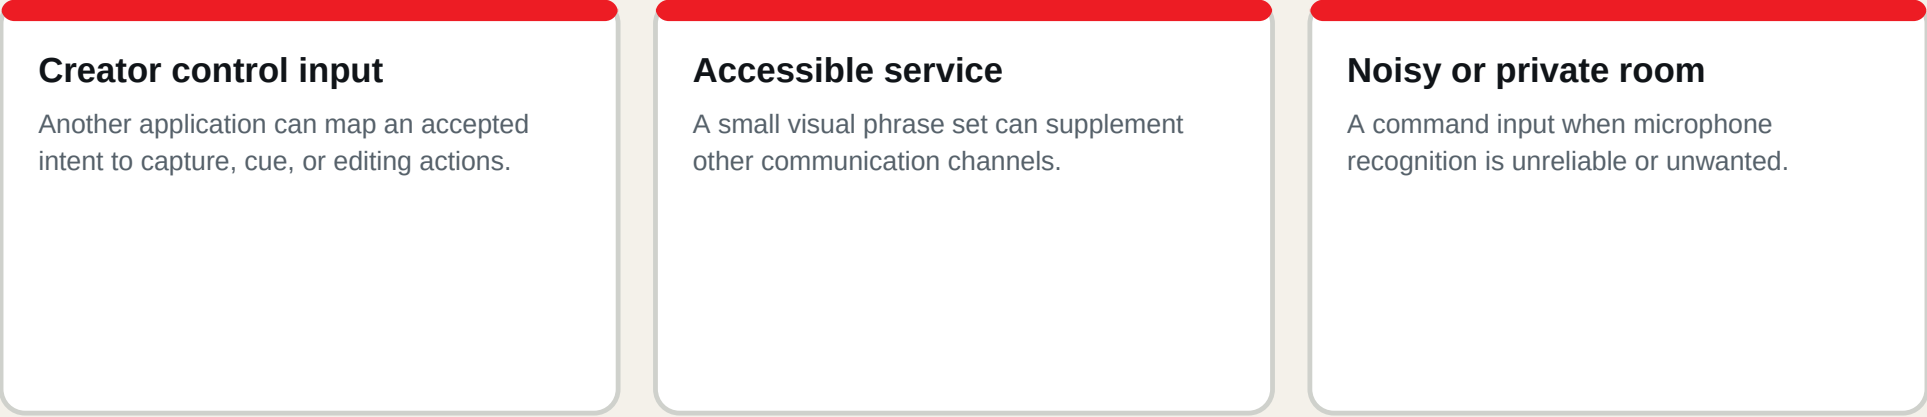
REGISTERED PHRASES

你好，请帮我打开灯 | 你吃饭了吗？

HOW IT WORKS



WHERE IT FITS



SAFETY RULE

Both calibrated gates must pass. Rejected clips cannot execute.

AMD RADEON + ROCM

CPU handles video and face preprocessing. ROCm PyTorch runs the learned fixed-phrase model on Radeon cuda:0.

EVIDENCE STATUS

ROCm execution and final evaluation remain to be recorded. Small-data smoke proves pipeline execution only.

SOURCE REPOSITORY